

assignment_3_mta

October 27, 2025

1 Assignment 3

1.0.1 Due 9/23. Do four of five.

1.

- Open the NHANES (or Ames prices or college completion datasets, if you prefer)
- Find two categorical variables of interest (there are 198, and short descriptions are given in the `nhanes_meta_17_18.csv` file). Investigate their missing values (you don't have to focus on missing values for this analysis like we did with police use of force, but always be aware of how dirty the data are)
- Compute a contingency table for your categorical X and Y
- Discuss any interesting patterns (or lack of one) that you observe

```
[41]: import os
import pandas as pd
import seaborn as sns

os.chdir("/Users/mta/Desktop/Understanding Uncertainty/MTA4KD/data")

data = pd.read_csv("nhanes_data_17_18.csv")
meta = pd.read_csv("nhanes_meta_17_18.csv")

print("Participants:", data.shape[0])
print("Variables:", data.shape[1])
print("Metadata variables:", meta.shape[0])

display(data.head())
display(meta[['Variable', 'Label', 'Target']].head(200))
```

Participants: 8366

Variables: 198

Metadata variables: 197

```
/var/folders/_x/p9ds2_7j4cg3964ybwvy6rh00000gn/T/ipykernel_83926/2540604510.py:7
: DtypeWarning: Columns (142) have mixed types. Specify dtype option on import
or set low_memory=False.
```

```
data = pd.read_csv("nhanes_data_17_18.csv")
```

```
SEQN GeneralHealthCondition EverBreastfedOrFedBreastmilk \
```

0	93703.0	NaN	1.0
1	93704.0	NaN	1.0
2	93705.0	Good	NaN
3	93706.0	Very good	NaN
4	93707.0	Good	NaN

	AgeStoppedBreastfeedingdays	AgeFirstFedFormuladays	\
0	273.0	1.0	
1	60.0	3.0	
2	NaN	NaN	
3	NaN	NaN	
4	NaN	NaN	

	AgeStoppedReceivingFormuladays	AgeStartedOtherFoodbeverage	\
0	365.0	152.0	
1	365.0	126.0	
2	NaN	NaN	
3	NaN	NaN	
4	NaN	NaN	

	AgeFirstFedMilkdays	TypeOfMilkFirstFedWholeMilk	TypeOfMilkFirstFed2Milk	\
0	365.0	Whole or regular milk	NaN	
1	365.0	Whole or regular milk	NaN	
2	NaN	NaN	NaN	
3	NaN	NaN	NaN	
4	NaN	NaN	NaN	

...	DaysSmokedCigsDuringPast30Days	AvgCigarettesdayDuringPast30Days	\
0	...	NaN	NaN
1	...	NaN	NaN
2	...	NaN	NaN
3	...	NaN	NaN
4	...	NaN	NaN

	TriedToQuitSmoking	TimesStoppedSmokingCigarettes	\
0	NaN	NaN	
1	NaN	NaN	
2	NaN	NaN	
3	NaN	NaN	
4	NaN	NaN	

	HowLongWereYouAbleToStopSmoking	UnitOfMeasureDayweekmonthyear_2_SMQ	\
0	NaN	NaN	
1	NaN	NaN	
2	NaN	NaN	
3	NaN	NaN	
4	NaN	NaN	

	CurrentSelfreportedHeightInches	CurrentSelfreportedWeightPounds	\
0	NaN	NaN	
1	NaN	NaN	
2	63.0	165.0	
3	68.0	145.0	
4	NaN	NaN	

	TriedToLoseWeightInPastYear	TimesLost10LbsOrMoreToLoseWeight
0	NaN	NaN
1	NaN	NaN
2	0.0	11 times or more
3	0.0	Never
4	NaN	NaN

[5 rows x 198 columns]

	Variable	Label	\
0	HSD010	General health condition	
1	DBQ010	Ever breastfed or fed breastmilk	
2	DBD030	Age stopped breastfeeding(days)	
3	DBD041	Age first fed formula(days)	
4	DBD050	Age stopped receiving formula(days)	
..	...	...	
192	SMQ852U	Unit of measure (day/week/month/year)	
193	WHD010	Current self-reported height (inches)	
194	WHD020	Current self-reported weight (pounds)	
195	WHQ070	Tried to lose weight in past year	
196	WHQ225	Times lost 10 lbs or more to lose weight	

	Target
0	Both males and females 12 YEARS -\n\n\t\t\t15...
1	Both males and females 0 YEARS -\n\n\t\t\t6 Y...
2	Both males and females 0 YEARS -\n\n\t\t\t6 Y...
3	Both males and females 0 YEARS -\n\n\t\t\t6 Y...
4	Both males and females 0 YEARS -\n\n\t\t\t6 Y...
..	...
192	Both males and females 18 YEARS -\n\n\t\t\t15...
193	Both males and females 16 YEARS -\n\n\t\t\t15...
194	Both males and females 16 YEARS -\n\n\t\t\t15...
195	Both males and females 16 YEARS -\n\n\t\t\t15...
196	Both males and females 16 YEARS -\n\n\t\t\t15...

[197 rows x 3 columns]

I'm choosing to analyze 1) Race/Hispanic origin w/ NH Asian (RIDRETH3) and 2) Doctor ever said you had arthritis (MCQ160A). These are both categorical variables.


```

--- DoctorEverSaidYouHadArthritis (MCQ160A) ---
Variable
EnglishText                                Label
MCQ160A Has a doctor or other health professional ever told {you/SP} that
{you/s/he} . . .had arthritis (ar-thry-tis)? Doctor ever said you had arthritis
DoctorEverSaidYouHadArthritis  Frequency
0                               0.0        3645
1                               NaN        3116
2                               1.0        1605

```

```

[45]: selected_vars = [col_race, col_arthritis]

missing_counts = pd.DataFrame({
    var: data[var].isna().value_counts()
    for var in selected_vars
}).T

missing_pct = pd.DataFrame({
    var: data[var].isna().value_counts(normalize=True) * 100
    for var in selected_vars
}).T

print("\nMissingness Counts")
print(missing_counts)

print("\nMissingness Percentages")
print(missing_pct.round(2))

```

Missingness Counts

	False	True
RacehispanicOriginWNhAsian	8366.0	NaN
DoctorEverSaidYouHadArthritis	5250.0	3116.0

Missingness Percentages

	False	True
RacehispanicOriginWNhAsian	100.00	NaN
DoctorEverSaidYouHadArthritis	62.75	37.25

```

[46]: contingency_counts = pd.crosstab(
    data[col_race],
    data[col_arthritis],
    margins=True,
    dropna=False
)

contingency_row_pct = pd.crosstab(
    data[col_race],

```

```

    data[col_arthritis],
    normalize="index",
    dropna=False
) * 100

contingency_col_pct = pd.crosstab(
    data[col_race],
    data[col_arthritis],
    normalize="columns",
    dropna=False
) * 100

print("\nContingency Table: Counts")
print(contingency_counts)

print("\nContingency Table: Row Percentages (within each race)")
print(contingency_row_pct.round(2))

print("\nContingency Table: Column Percentages (within arthritis response)")
print(contingency_col_pct.round(2))

```

Contingency Table: Counts

DoctorEverSaidYouHadArthritis	0.0	1.0	NaN	All
RacehispanicOriginWNhAsian				
Mexican American	544	152	533	1229
Non-Hispanic Asian	636	121	326	1083
Non-Hispanic Black	841	395	713	1949
Non-Hispanic White	1097	707	988	2792
Other Hispanic	359	133	246	738
Other Race - Including Multi-Racial	168	97	310	575
All	3645	1605	0	8366

Contingency Table: Row Percentages (within each race)

DoctorEverSaidYouHadArthritis	0.0	1.0	NaN
RacehispanicOriginWNhAsian			
Mexican American	44.26	12.37	43.37
Non-Hispanic Asian	58.73	11.17	30.10
Non-Hispanic Black	43.15	20.27	36.58
Non-Hispanic White	39.29	25.32	35.39
Other Hispanic	48.64	18.02	33.33
Other Race - Including Multi-Racial	29.22	16.87	53.91

Contingency Table: Column Percentages (within arthritis response)

DoctorEverSaidYouHadArthritis	0.0	1.0	NaN
RacehispanicOriginWNhAsian			
Mexican American	14.92	9.47	17.11
Non-Hispanic Asian	17.45	7.54	10.46

Non-Hispanic Black	23.07	24.61	22.88
Non-Hispanic White	30.10	44.05	31.71
Other Hispanic	9.85	8.29	7.89
Other Race - Including Multi-Racial	4.61	6.04	9.95

Looking at the contingency table, two things jump out. First, the row percentages: within each race, Non-Hispanic Whites have the highest reported arthritis rate (about 25%), followed by Non-Hispanic Blacks (20%). Mexican Americans and Asians are lower, closer to 12%. So prevalence clearly isn't even across groups. Second, the column percentages: when you look at the arthritis "Yes" pool, it's dominated by Whites (44%) and Blacks (25%), meaning those two groups make up the bulk of reported arthritis cases in this dataset. The big caveat is missingness — it's huge, 30–50% depending on the group — so you can't treat these as clean prevalence estimates. Still, the directional pattern is obvious: arthritis is more common among Whites and Blacks, less so among Mexican Americans and Asians.

2.

- Open the NHANES dataset
- Find a categorical and numeric variable of interest (there are 198, and short descriptions are given in the `nhanes_meta_17_18.csv` file). Investigate their missing values (you don't have to focus on missing values for this analysis, but always be aware of them)
- Make descriptive tables and grouped kernel density plots to represent the variation in your numeric Y conditional on your categorical X
- Discuss any interesting patterns (or lack of one) that you observe

I'm choosing to analyze 1) Weight KG (BMXWT) and 2) Doctor ever said you had arthritis (MCQ160A).

```
[47]: col_cat = "DoctorEverSaidYouHadArthritis"
      col_num = "WeightKg"

      for col in [col_cat, col_num]:
          missing_pct = data[col].isna().mean() * 100
          print(f"{col} - Missing: {missing_pct:.2f}%")

      desc_stats = data.groupby(col_cat)[col_num].describe()
      print("\nDescriptive Statistics (Weight by Arthritis status):")
      print(desc_stats)

      plt.figure(figsize=(8,5))
      sns.kdeplot(
          data=data,
          x=col_num,
          hue=col_cat,
          common_norm=False,
          fill=True,
          alpha=0.4
      )
```

```
plt.title("Distribution of Weight (kg) by Arthritis Status")
plt.xlabel("Weight (kg)")
plt.ylabel("Density")
plt.show()



plt.figure(figsize=(6,4))
sns.boxplot(
    data=data,
    x=col_cat,
    y=col_num
)
plt.title("Boxplot of Weight (kg) by Arthritis Status")
plt.xlabel("Arthritis Status (0=No, 1=Yes)")
plt.ylabel("Weight (kg)")
plt.show()
```

DoctorEverSaidYouHadArthritis - Missing: 37.25%

WeightKg - Missing: 1.48%

Descriptive Statistics (Weight by Arthritis status):

	count	mean	std	min	25%	50%	\
DoctorEverSaidYouHadArthritis							
0.0	3600.0	81.367611	22.227085	32.6	66.1	77.7	
1.0	1573.0	86.150223	23.990346	39.3	69.3	82.6	
	75%	max					
DoctorEverSaidYouHadArthritis							
0.0	92.7	242.6					
1.0	98.8	219.6					



Looking at the stats and plots, there's a clear weight gap between people who reported arthritis and those who didn't. Average weight is about 86 kg in the arthritis group vs 81 kg in the no-arthritis group. The kernel density plot shows the arthritis group's distribution shifted slightly to the right — heavier on average, with a longer right tail. The boxplot tells the same story: higher median and more extreme outliers among those with arthritis. Missingness is low here (~1.5% for weight, ~37% for arthritis status), so this pattern isn't just noise. The takeaway is straightforward: in this sample, higher weight tracks with higher arthritis prevalence.

3. We showed that the mean and median could be discovered by minimizing various kinds of loss functions; this is what machine learning is. To make a prediction $\hat{y}(z)$ of Y when $X = z$, minimize the mean squared error:

$$MSE(\hat{y}(z)) = \frac{1}{N} \sum_{i=1}^N \{y_i - \hat{y}(z)\}^2 \frac{1}{h} k\left(\frac{z - x_i}{h}\right)$$

Show that the solution to this problem is the LCLS/Naradaya-Watson estimator.

$$MSE(\hat{y}(z)) = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}(z))^2 w_i(z), \quad w_i(z) = \frac{1}{h} k\left(\frac{z - x_i}{h}\right).$$

Step 1: Differentiate wrt $\hat{y}(z)$

$$\frac{\partial}{\partial \hat{y}(z)} MSE(\hat{y}(z)) = \frac{1}{N} \sum_{i=1}^N 2(\hat{y}(z) - y_i) w_i(z) = \frac{2}{N} \left[\hat{y}(z) \sum_{i=1}^N w_i(z) - \sum_{i=1}^N w_i(z) y_i \right].$$

Step 2: Set equal to zero and solve

$$\begin{aligned} \hat{y}(z) \sum_{i=1}^N w_i(z) &= \sum_{i=1}^N w_i(z) y_i, \\ \hat{y}(z) &= \frac{\sum_{i=1}^N w_i(z) y_i}{\sum_{i=1}^N w_i(z)}. \end{aligned}$$

Step 3: Plug back in kernel weights

Since

$$w_i(z) = \frac{1}{h} k\left(\frac{z - x_i}{h}\right),$$

we get:

$$\hat{y}(z) = \frac{\sum_{i=1}^N k\left(\frac{z - x_i}{h}\right) y_i}{\sum_{i=1}^N k\left(\frac{z - x_i}{h}\right)}.$$

This is the **Nadaraya–Watson estimator** - a weighted average of the y_i 's, with weights coming from the kernel.

- Small h : sharp, twitchy, low bias, high variance. Like trading tick-by-tick — overfit and volatile.
- Large h : smooth, flat, high bias, low variance. Like averaging across decades — stable, but you miss the details.

The *estimator* is just balancing how much noise we want to eat versus how much structure we're willing to miss.

4.

- Write a class or set of functions that implement the LCLS/Nadaraya-Watson estimator, using the Silverman plug-in estimate for the conditioning variable X as the bandwidth.
- From one of the course data sets, find two numeric variables of interest, analyze their relationship with the the LCLS/Nadaraya-Watson estimator, and discuss your results.

```
[48]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from scipy.stats import norm

col_x = "BodyMassIndexKgm2"
col_y = "TotalCholesterolMgdl"

subset = data[[col_x, col_y]].dropna()
X = subset[col_x].values
Y = subset[col_y].values

n = len(X)
h = 1.06 * np.std(X) * n ** (-1/5)


def nw_estimator(x0, X, Y, h):
    weights = norm.pdf((x0 - X) / h)
    return np.sum(weights * Y) / np.sum(weights)

x_grid = np.linspace(X.min(), X.max(), 200)
y_hat = np.array([nw_estimator(x0, X, Y, h) for x0 in x_grid])

plt.figure(figsize=(8,6))
plt.scatter(X, Y, s=10, alpha=0.3, label="Observed data")
plt.plot(x_grid, y_hat, color="red", lw=2, label="Nadaraya-Watson estimate")
plt.xlabel("Body Mass Index (kg/m^2)")
plt.ylabel("Total Cholesterol (mg/dL)")
plt.title("NW Estimator: BMI vs Total Cholesterol")
plt.legend()
```

```
plt.show()

print("Descriptive stats:")
print(subset.describe())
```



Descriptive stats:

	BodyMassIndexKgm2	TotalCholesterolMgdl
count	6641.000000	6641.000000
mean	27.833444	179.982232
std	7.913572	40.647648
min	12.600000	76.000000
25%	22.300000	150.000000
50%	26.900000	176.000000
75%	32.100000	204.000000
max	86.200000	446.000000

The scatterplot shows the relationship between BMI and Total Cholesterol. The raw data are highly variable, but the Nadaraya-Watson estimator provides a smoothed curve that highlights the average trend. Cholesterol levels tend to increase slightly as BMI rises from low to mid-range, then flatten, with some increase again at higher BMI values.

The descriptive statistics confirm the wide spread in both variables: BMI (std = 7.9) and cholesterol (std = 49.6). Missingness is relatively modest for BMI (4%) but higher for cholesterol (19%).

Overall, the estimator suggests a weak positive association between BMI and cholesterol, though with considerable noise. The pattern indicates that individuals with higher BMI often show higher cholesterol on average, but the relationship is far from deterministic, and there is substantial variation across the range of BMI values.

5.

- In any of the available data sets, investigate the relationships between pairs of variables (X, Y) with a scatterplot and CEF (for example, price on area)
- Is this relationship plausibly causal, or are there missing variables that might explain at least part of the relationship between your variables? These can be “conceptual” rather than “practical”; for example, ‘talent’ or ‘grit’ probably explain education outcomes, but are almost impossible to measure. We are asking whether there are hypothetical **threats to causal identification** of the effect of X on Y .
- Explain how, regardless of the threat to causal identification, you can still use your model to predict Y given X , as long as you don’t intervene in the system to control the outcome